

final

December 20, 2025

1 initialize model

```
[183]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from xgboost import XGBRegressor
from lightgbm import LGBMRegressor
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.metrics import (
    mean_absolute_error,
    mean_squared_error,
    r2_score,
)
```

```
[184]: df = pd.read_csv('data.csv')
df.rename(columns={ df.columns[0]: "date" }, inplace = True)
```

2 XGBoost

```
[185]: base_features = [
    '^VIX',
    '^GSPC',
    'CCC_10Y_Spread',
    'Unemployment',
    'Inflation',
    'GDP_Growth',
    'Corporate_Leverage',
    'YieldCurveSlope',
    'IndPro_Growth',
    'CP_Growth',
]

# Assume df is already loaded externally
```

```

df = df.sort_values("date")

# Convenience alias
df["spread"] = df["CCC_10Y_Spread"]

# # Directional features on spread (using levels)
# df["spread_ret_30"] = df["spread"].diff(30)
# df["spread_ret_60"] = df["spread"].diff(60)
# df["spread_ret_180"] = df["spread"].diff(180)
# df["spread_ret_360"] = df["spread"].diff(365)

df["spread_vol_30"] = df["spread"].diff().rolling(30).std()
df["spread_vol_60"] = df["spread"].diff().rolling(60).std()
df["spread_vol_180"] = df["spread"].diff().rolling(180).std()
df["spread_vol_360"] = df["spread"].diff().rolling(365).std()

# First-difference (change) of base macro features
for col in base_features:
    df[col + "_week"] = df[col].diff(7)
    df[col + "_month"] = df[col].diff(30)
    df[col + "_twomonth"] = df[col].diff(60)
    df[col + "_halfyear"] = df[col].diff(180)
    df[col + "_year"] = df[col].diff(365)

# Drop rows with NaNs from shifting/diff/rolling
df = df.dropna().copy()

train_end = '2016-12-31'
val_end = '2021-12-31'

train = df[df['date'] <= train_end]
val = df[(df['date'] > train_end) & (df['date'] <= val_end)]
test = df[df['date'] > val_end]

# Numeric features: original + engineered
feature_cols_num = (
    base_features +
    [
        # "spread_ret_30", "spread_ret_60",
        # "spread_ret_180", "spread_ret_360",
        "spread_vol_30", "spread_vol_60",
        "spread_vol_180", "spread_vol_360",]
    + [col + "_twomonth" for col in base_features] +
    [col + "_month" for col in base_features] +

```

```

        [col + "_halfyear" for col in base_features] +
        [col + "_year" for col in base_features]
    )

    # If you later add categoricals, put names here
    feature_cols_cat = [] # e.g. ['rating', 'sector']

    X_train = train[feature_cols_num]
    X_val    = val[feature_cols_num]
    X_test   = test[feature_cols_num]

    X_train_lgb = X_train
    X_val_lgb   = X_val
    X_test_lgb  = X_test

    # Target is the spread LEVEL
    y_train = train["spread"]
    y_val    = val["spread"]
    y_test   = test["spread"]

    preprocess = ColumnTransformer(
        transformers=[
            ('num', StandardScaler(), feature_cols_num),
            ('cat', OneHotEncoder(handle_unknown='ignore'), feature_cols_cat),
        ],
        remainder='drop'
    )

    preprocess.fit(X_train)

    X_train_t = preprocess.transform(X_train)
    X_val_t    = preprocess.transform(X_val)
    X_test_t   = preprocess.transform(X_test)

```

```

[186]: xgb_model = XGBRegressor(
        n_estimators=3000,
        max_depth=0,
        learning_rate=0.02,
        # subsample=0.8,
        # colsample_bytree=0.8,
        # min_child_weight=5,
        # reg_lambda=2.0,
        # reg_alpha=1.0,
        objective='reg:squarederror',

```

```

        random_state=42,
    )

    xgb_model.fit(
        X_train_t, y_train,
        eval_set=[(X_val_t, y_val)],
        verbose=False
    )

```

```

[186]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                    colsample_bylevel=None, colsample_bynode=None,
                    colsample_bytree=None, device=None, early_stopping_rounds=None,
                    enable_categorical=False, eval_metric=None, feature_types=None,
                    feature_weights=None, gamma=None, grow_policy=None,
                    importance_type=None, interaction_constraints=None,
                    learning_rate=0.02, max_bin=None, max_cat_threshold=None,
                    max_cat_to_onehot=None, max_delta_step=None, max_depth=0,
                    max_leaves=None, min_child_weight=None, missing=nan,
                    monotone_constraints=None, multi_strategy=None, n_estimators=3000,
                    n_jobs=None, num_parallel_tree=None, ...)

```

```

[187]: lgbm_model = LGBMRegressor(
    n_estimators=3000,
    learning_rate=0.02,
    num_leaves=31,
    max_depth=-1, # let num_leaves control tree size
    subsample=0.8,
    colsample_bytree=0.8,
    objective='regression',
    random_state=42, # if your version accepts n_jobs, you can add: n_jobs=-1
)

lgbm_model.fit(
    X_train_lgb, y_train,
    eval_set=[(X_val_lgb, y_val)],
    eval_metric='rmse',
)

```

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000505 seconds.
 You can set `force\_col\_wise=true` to remove the overhead.
 [LightGBM] [Info] Total Bins 13769
 [LightGBM] [Info] Number of data points in the train set: 6941, number of used features: 54
 [LightGBM] [Info] Start training from score 0.114654

```

[187]: LGBMRegressor(colsample_bytree=0.8, learning_rate=0.02, n_estimators=3000,
                    objective='regression', random_state=42, subsample=0.8)

```

3 Metrics

3.1 XGBoost

```
[188]: def metrics_with_direction(y_true, y_pred):
    y_true = np.array(y_true)
    y_pred = np.array(y_pred)

    mae = mean_absolute_error(y_true, y_pred)
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    r2 = r2_score(y_true, y_pred)

    # MAPE on levels
    denom = np.where(y_true == 0, 1e-8, y_true)
    mape = np.mean(np.abs((y_true - y_pred) / denom)) * 100

    # Directional Accuracy based on level changes
    true_dir = np.sign(y_true[1:] - y_true[:-1])
    pred_dir = np.sign(y_pred[1:] - y_pred[:-1])
    da = np.mean(true_dir == pred_dir)

    return mae, rmse, r2, mape, da

[189]: y_pred_train_x = xgb_model.predict(X_train_t)
    y_pred_val_x = xgb_model.predict(X_val_t)

    train_mae, train_rmse, train_r2, train_mape, train_da = metrics_with_direction(
        y_train.values, y_pred_train_x
    )
    val_mae, val_rmse, val_r2, val_mape, val_da = metrics_with_direction(
        y_val.values, y_pred_val_x
    )

    print("=== TRAIN METRICS (level) ===")
    print(f"MAE : {train_mae:.4f}")
    print(f"RMSE: {train_rmse:.6f}")
    print(f"R2 : {train_r2:.4f}")
    print(f"MAPE: {train_mape:.2f}%")
    print(f"DA : {train_da:.4f}")

    print("\n=== VALIDATION METRICS (level) ===")
    print(f"MAE : {val_mae:.4f}")
    print(f"RMSE: {val_rmse:.6f}")
    print(f"R2 : {val_r2:.4f}")
    print(f"MAPE: {val_mape:.2f}%")
    print(f"DA : {val_da:.4f}")

    print("=== TRAIN METRICS (level) ===")
```

MAE : 0.0002
RMSE: 0.000250
R² : 1.0000
MAPE: 0.15%
DA : 0.6807

=== VALIDATION METRICS (level) ===

MAE : 0.0003
RMSE: 0.000859
R² : 0.9989
MAPE: 0.27%
DA : 0.6449

```
[190]: y_pred_train = lgbm_model.predict(X_train)
y_pred_val = lgbm_model.predict(X_val)

train_mae, train_rmse, train_r2, train_mape, train_da = metrics_with_direction(
    y_train.values, y_pred_train
)
val_mae, val_rmse, val_r2, val_mape, val_da = metrics_with_direction(
    y_val.values, y_pred_val
)

print("=== TRAIN METRICS (level) ===")
print(f"MAE : {train_mae:.4f}")
print(f"RMSE: {train_rmse:.6f}")
print(f"R2 : {train_r2:.4f}")
print(f"MAPE: {train_mape:.2f}%")
print(f"DA : {train_da:.4f}")

print("\n=== VALIDATION METRICS (level) ===")
print(f"MAE : {val_mae:.4f}")
print(f"RMSE: {val_rmse:.6f}")
print(f"R2 : {val_r2:.4f}")
print(f"MAPE: {val_mape:.2f}%")
print(f"DA : {val_da:.4f}")
```

=== TRAIN METRICS (level) ===

MAE : 0.0001
RMSE: 0.000161
R² : 1.0000
MAPE: 0.10%
DA : 0.9056

=== VALIDATION METRICS (level) ===

MAE : 0.0035
RMSE: 0.005482
R² : 0.9564

MAPE: 3.87%
DA : 0.7808

4 Plots

```
[191]: plt.figure(figsize=(7, 7))
plt.scatter(y_val, y_pred_val_x, alpha=0.5)

min_val = min(y_val.min(), y_pred_val_x.min())
max_val = max(y_val.max(), y_pred_val_x.max())
plt.plot([min_val, max_val], [min_val, max_val], linestyle='--')

plt.xlabel("Actual Spread")
plt.ylabel("Predicted Spread")
plt.title("Actual vs Predicted (Validation, Levels)")
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Time series: Actual vs Predicted over time ---
val_df = pd.DataFrame(
    {
        "Actual": y_val.values,
        "Predicted": y_pred_val_x,
    },
    index=val["date"]
)

plt.figure(figsize=(12, 5))
plt.plot(val_df.index, val_df["Actual"], label="Actual")
plt.plot(val_df.index, val_df["Predicted"], label="Predicted")
plt.xlabel("Time")
plt.ylabel("Spread Level")
plt.title("Actual vs Predicted Spread over Time (Validation)")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Feature Importances ---
importances = xgb_model.feature_importances_

# importances = lgbm_model.feature_importances_

# Indices of non-zero importances
nz_idx = np.where(importances > 0.001)[0]
```


```


# Sorted indices (only non-zero)
sorted_nz_idx = nz_idx[np.argsort(importances[nz_idx])]


# Corresponding feature names
nz_features = np.array(feature_cols_num)[sorted_nz_idx]
nz_importances = importances[sorted_nz_idx]

plt.figure(figsize=(8, 8))
plt.barh(range(len(sorted_nz_idx)), nz_importances)
plt.yticks(range(len(sorted_nz_idx)), nz_features)
plt.xlabel("Importance")
plt.title("XGBoost Non-Zero Feature Importances")
plt.tight_layout()
plt.show()

```







```
[192]: plt.figure(figsize=(7, 7))
plt.scatter(y_val, y_pred_val, alpha=0.5)

min_val = min(y_val.min(), y_pred_val.min())
max_val = max(y_val.max(), y_pred_val.max())
plt.plot([min_val, max_val], [min_val, max_val], linestyle='--')

plt.xlabel("Actual Spread")
plt.ylabel("Predicted Spread")
plt.title("Actual vs Predicted (Validation, Levels)")
plt.grid(True)
plt.tight_layout()
plt.show()
```

```

# --- Time series: Actual vs Predicted over time ---
val_df = pd.DataFrame(
    {
        "Actual": y_val.values,
        "Predicted": y_pred_val,
    },
    index=val["date"]
)

plt.figure(figsize=(12, 5))
plt.plot(val_df.index, val_df["Actual"], label="Actual")
plt.plot(val_df.index, val_df["Predicted"], label="Predicted")
plt.xlabel("Time")
plt.ylabel("Spread Level")
plt.title("Actual vs Predicted Spread over Time (Validation)")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Feature Importances ---
# importances = xgb_model.feature_importances_

importances = lgbm_model.feature_importances_

# Indices of non-zero importances
nz_idx = np.where(importances > 0.001)[0]

# Sorted indices (only non-zero)
sorted_nz_idx = nz_idx[np.argsort(importances[nz_idx])]

# Corresponding feature names
nz_features = np.array(feature_cols_num)[sorted_nz_idx]
nz_importances = importances[sorted_nz_idx]

plt.figure(figsize=(8, 8))
plt.barh(range(len(sorted_nz_idx)), nz_importances)
plt.yticks(range(len(sorted_nz_idx)), nz_features)
plt.xlabel("Importance")
plt.title("LightGBM Non-Zero Feature Importances")
plt.tight_layout()
plt.show()

```

